

Push is Faster



using `std::cpp` 2025

Joaquín M López Muñoz <joaquin.lopezmunoz@gmail.com>

Madrid, March 2025



```
for(auto x: rng) {  
    if(x%2 == 0) std::cout << x*3 << " ";  
}
```


```
for(auto x: rng) {  
    if(x%2 == 0) std::cout << x*3 << " ";  
}
```






```
auto is_even = [](auto x) { return x%2 == 0; };  
auto x3 =      [](auto x) { return x*3; };  
auto dst =      [](auto x) { std::cout << x << " "; };  
  
for(auto x: rng) {  
    if(is_even(x)) dst(x3(x));  
}
```

```
for(auto x: rng) {  
    if(x%2 == 0) std::cout << x*3 << " ";  
}
```

```
auto is_even = [](auto x) { return x%2 == 0; };  
auto x3 =      [](auto x) { return x*3; };  
auto dst =     [](auto x) { std::cout << x << " "; };  
  
for(auto x: rng) {  
    if(is_even(x)) dst(x3(x));  
}
```

```
using namespace std::views;  
  
for (auto x: rng | filter(is_even) | transform(x3)) {  
    dst(x);  
}
```





Push or pull



Push passes values

```
for(auto x: rng) {  
    if(is_even(x)) dst(x3(x));  
}
```



```
auto filter = [](auto pred, auto next) {  
    return [=](auto x) { if(pred(x)) next(x); };  
};  
auto transform = [](auto f, auto next) {  
    return [=](auto x) { next(f(x)); };  
};  
  
for(auto x: rng) {  
    filter(is_even, transform(x3, dst))(x);  
}
```

Pull retrieves values


```
for (auto x:  
    rng | filter(is_even) | transform(x3)) {  
    dst(x);  
}
```



```
for (auto x: transform(filter(rng, is_even), x3)) {  
    dst(x);  
}
```

```
for(auto x: rng) {  
    if(is_even(x)) dst(x3(x));  
}
```

Push: control point at the beginning



Pull: control point at the end

```
for (auto x: rng | filter(is_even) | transform(x3)) {  
    dst(x);  
}
```




Push

- Reactive programming
- Continuation-passing style
- Senders/receivers (in a sense)


Pull

- STL iterators
- C++ ranges
- SQL (in a sense)

- All fan-in nodes must precede the control point
 - All fan-out nodes must follow the control point
- ⇒ Push can't do DAGs with fan-in
- ⇒ Pull can't do DAGs with fan-out
- ⇒ Some DAGs are neither pushable nor pullable




```
for (auto x:
    concat(rng1 | filter(is_even), rng2) |
    transform(x3)) {
    dst(x);
}
```



```
for(auto x: rng) {
    if(is_even(x)) {
        dst1(x);
        dst2(x3(x));
    }
}
```

↓

```
for(auto x: rng) {
    filter(is_even,
        tee(dst1, transform(x3, dst2)))(x);
}
```




```
std::vector<int> rng;  
int res;  
  
auto is_even = [](auto x) { return x%2 == 0; };  
auto x3 =      [](auto x) { return x*3; };  
void dst(int);  
  
void push()  
{  
    for(auto x: rng) {  
        if(is_even(x)) dst(x3(x));  
    }  
}  
  
void pull()  
{  
    for (auto x:  
        rng | filter(is_even) | transform(x3)) {  
        dst(x);  
    }  
}
```

<https://godbolt.org/z/o4KTKef59>

```

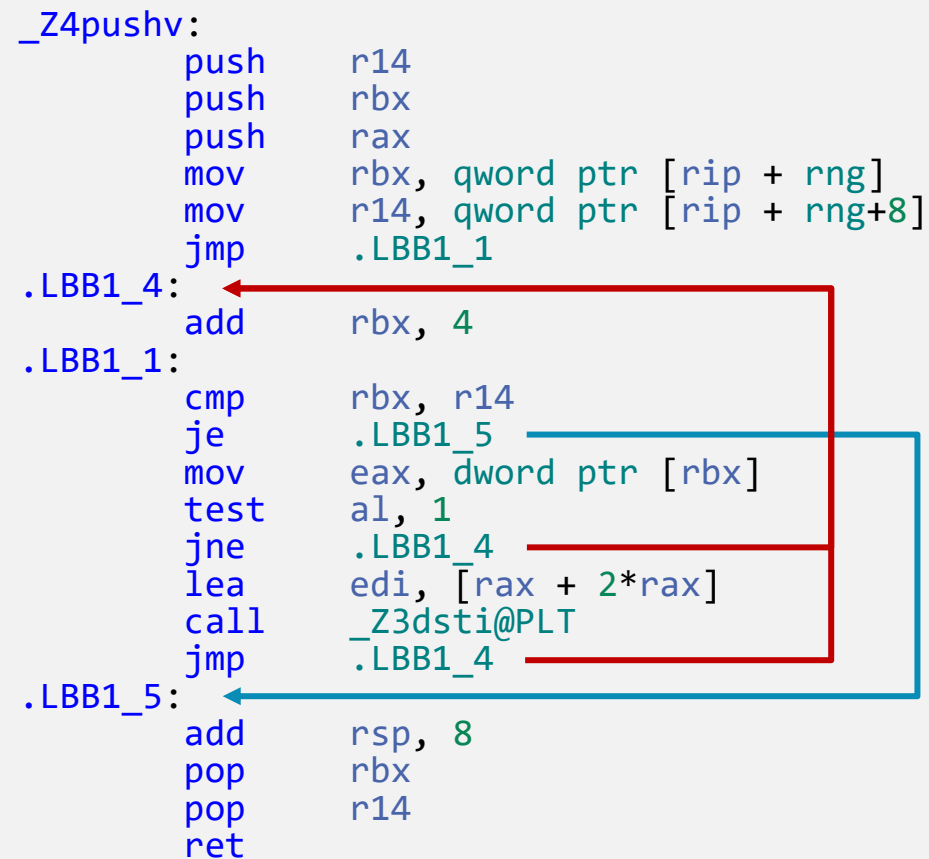
std::vector<int> rng;
int res;

auto is_even = [](auto x) { return x%2 == 0; };
auto x3 =      [](auto x) { return x*3; };
void dst(int);

void push()
{
    for(auto x: rng) {
        if(is_even(x)) dst(x3(x));
    }
}

void pull()
{
    for (auto x:
        rng | filter(is_even) | transform(x3)) {
        dst(x);
    }
}

```



```

_Z4pushv:
    push    r14
    push    rbx
    push    rax
    mov     rbx, qword ptr [rip + rng]
    mov     r14, qword ptr [rip + rng+8]
    jmp     .LBB1_1
.LBB1_4:
    add     rbx, 4
.LBB1_1:
    cmp     rbx, r14
    je      .LBB1_5
    mov     eax, dword ptr [rbx]
    test    al, 1
    jne     .LBB1_4
    lea     edi, [rax + 2*rax]
    call    _Z3dsti@PLT
    jmp     .LBB1_4
.LBB1_5:
    add     rsp, 8
    pop     rbx
    pop     r14
    ret

```

```

std::vector<int> rng;
int res;

auto is_even = [](auto x) { return x%2 == 0; };
auto x3 =      [](auto x) { return x*3; };
void dst(int);

void push()
{
    for(auto x: rng) {
        if(is_even(x)) dst(x3(x));
    }
}

void pull()
{
    for (auto x:
        rng | filter(is_even) | transform(x3)) {
        dst(x);
    }
}

```

```

_Z4pullv:
    push    r14
    push    rbx
    push    rax
    mov     rbx, qword ptr [rip + rng]
    mov     r14, qword ptr [rip + rng+8]
    cmp     rbx, r14
    je      .LBB2_4

.LBB2_1:
    test    byte ptr [rbx], 1
    je      .LBB2_4
    add     rbx, 4
    cmp     rbx, r14
    jne     .LBB2_1
    jmp     .LBB2_3

.LBB2_5:
    mov     eax, dword ptr [rbx]
    lea     edi, [rax + 2*rax]
    call    _Z3dsti@PLT
    add     rbx, 4
    mov     rax, qword ptr [rip + rng+8]
    cmp     rbx, rax
    je      .LBB2_4

.LBB2_7:
    test    byte ptr [rbx], 1
    je      .LBB2_4
    add     rbx, 4
    cmp     rbx, rax
    jne     .LBB2_7

.LBB2_4:
    cmp     rbx, r14
    jne     .LBB2_5

.LBB2_3:
    add     rsp, 8
    pop     rbx
    pop     r14
    ret

```



```
auto v = filter(rng, is_even);  
auto first = v.begin();  
auto last = v.end();  
  
while(first != last)  
{  
    dst(*first);  
    ++first;  
}
```

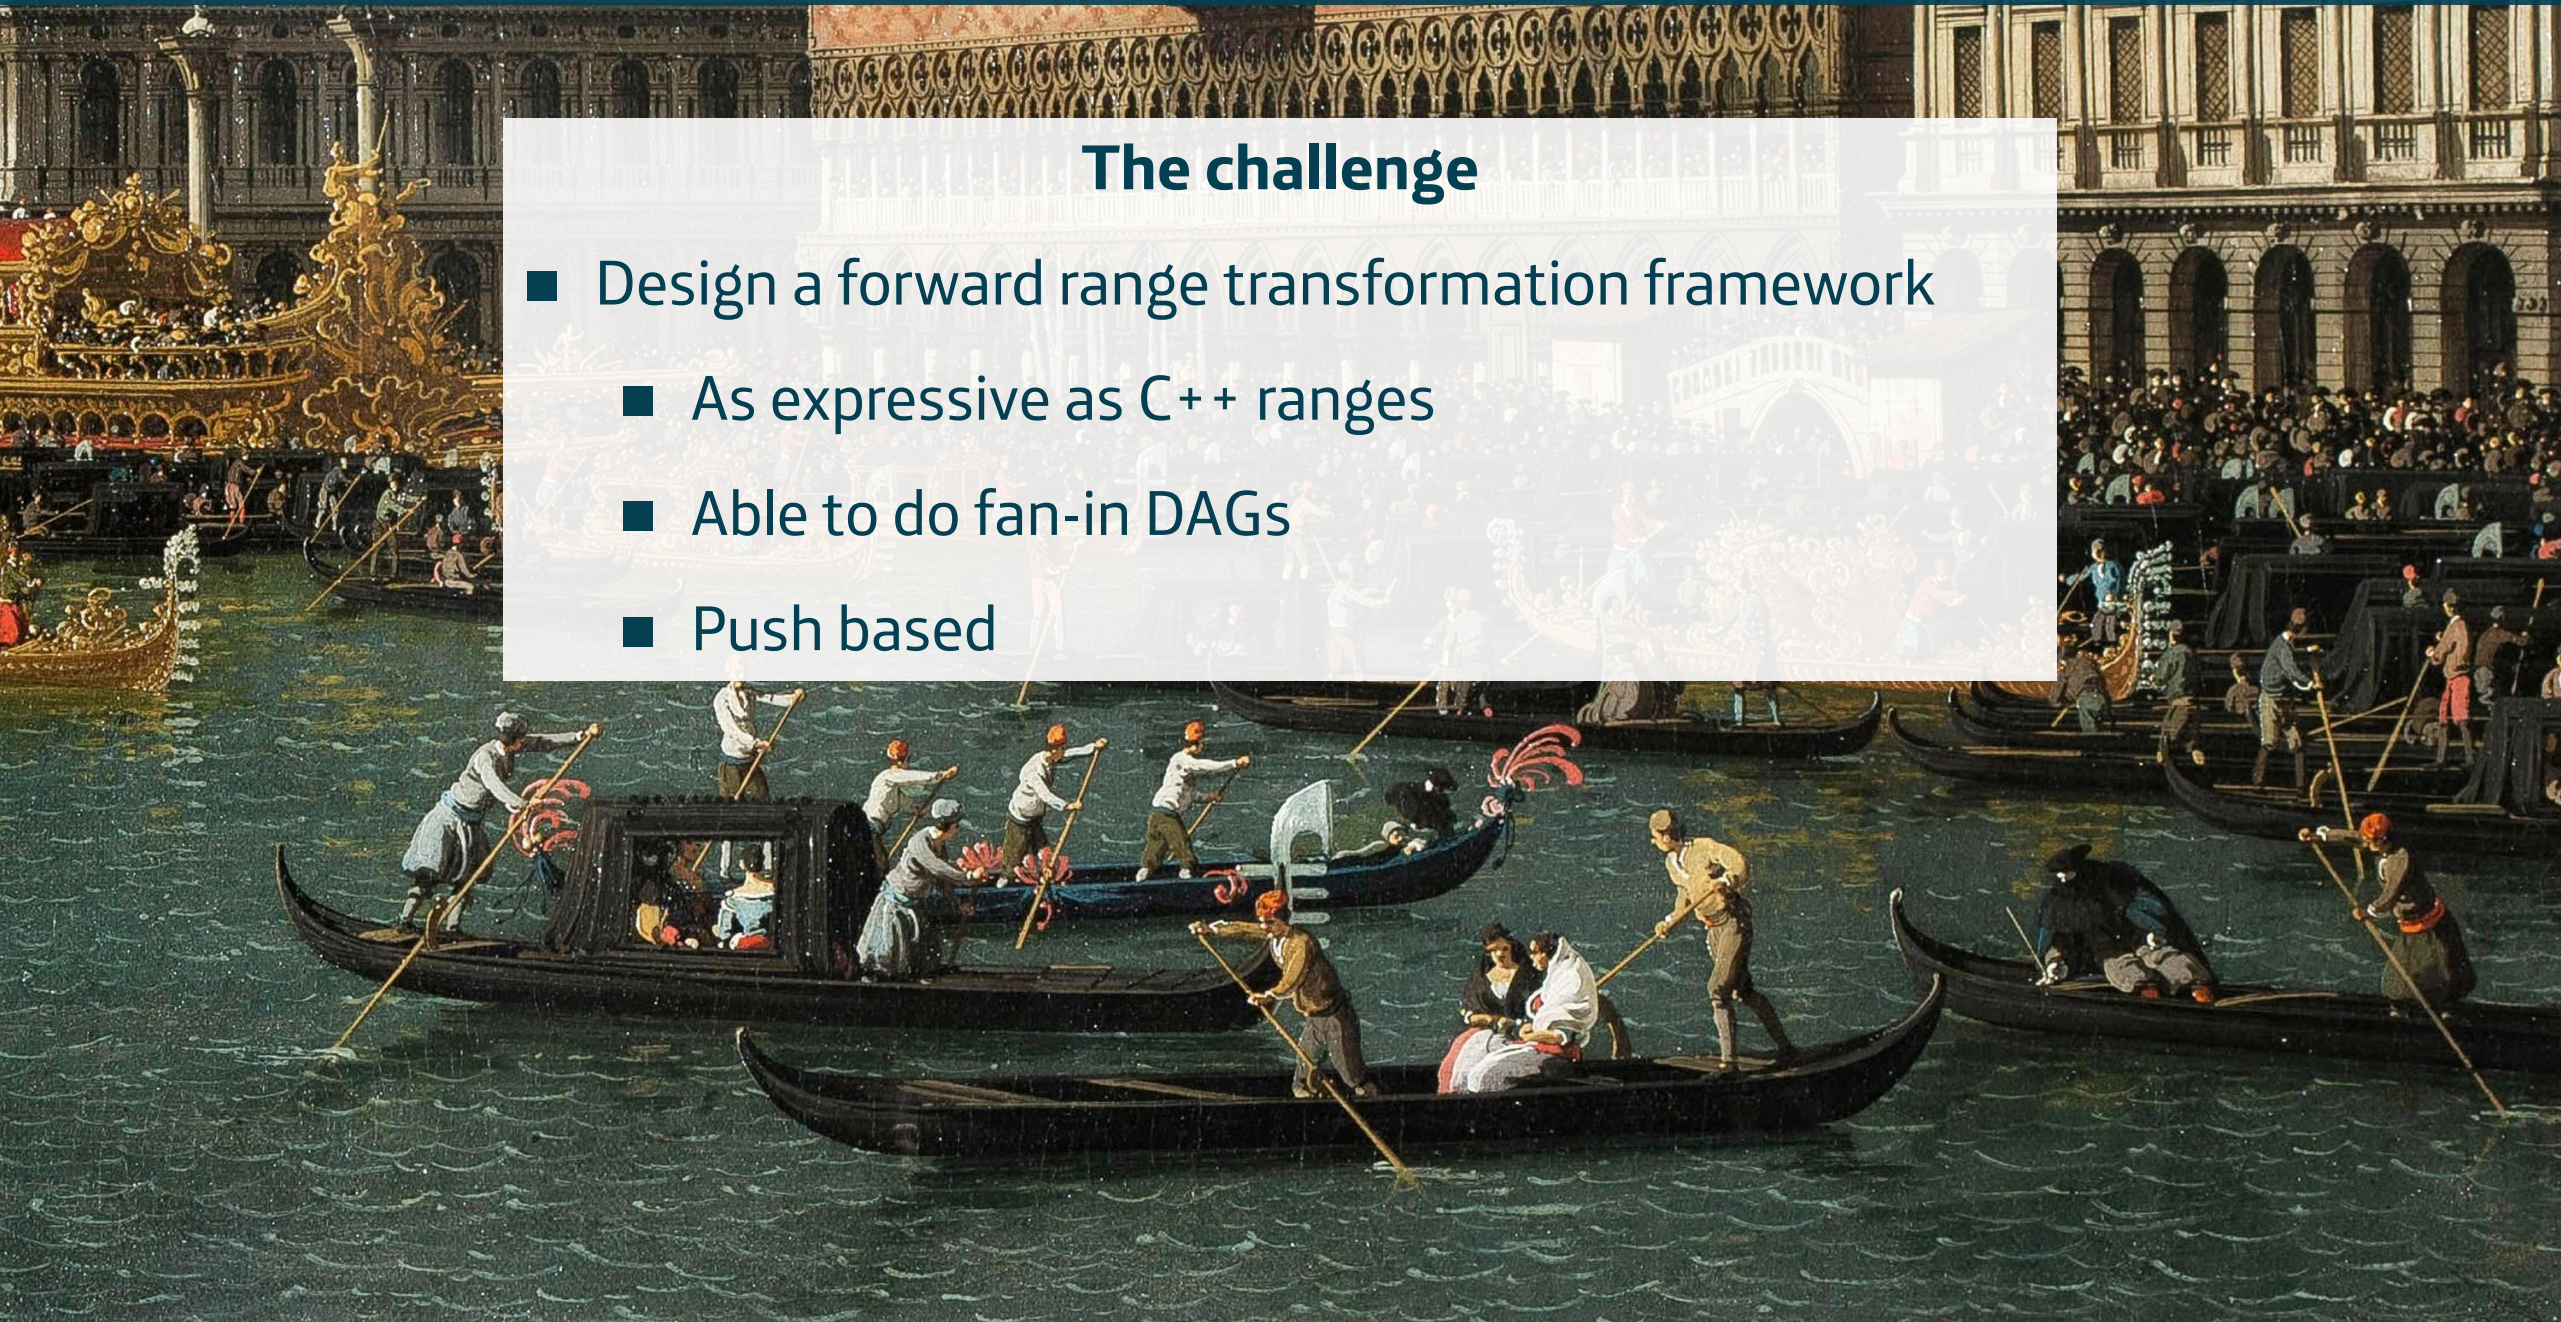
- Transformation is not 1:1 → double end-check
- **concat**, **join**, etc. have problems of their own
 - Context must be restored on each iteration
 - Fat iterators
- Compilers don't always see through all this
 - Much better at optimizing argument passing → context kept on call stack
 - Inversion of control → push

Transrangers



The challenge


- Design a forward range transformation framework
 - As expressive as C++ ranges
 - Able to do fan-in DAGs
 - Push based




```
template<typename Cursor>
concept cursor = requires(Cursor p)
{
    { *p } -> can_reference;
};

template<typename Dst, typename Cursor>
concept destination_function = requires(Dst dst, Cursor p)
{
    { dst(p) } -> std::convertible_to<bool>;
};

template<typename Ranger, typename Dst>
concept ranger = requires(Ranger rgr, Dst dst)
{
    typename Ranger::cursor;
    requires cursor<typename Ranger::cursor>;
    requires destination_function<Dst, typename Ranger::cursor>;
    { rgr(dst) } -> std::convertible_to<bool>;
};
```




- Traverses a range
- Passes cursors to `dst` while it returns `true`
- Returns `false` if there are still elements to process
- May process the range with several destination functions

```
#include <transrangers.hpp>

using namespace transrangers;

auto rgr = all(rng); // plain ranger over range rng
auto b = rgr([n = 10](auto p) mutable { // print 10 elements
    std::cout<< *p << " ";
    return --n != 0;
});
if(!b) { // still elements in the range
    rgr([](auto p) { // print remaining in brackets
        std::cout<< "[" << *p << "]" ";
        return true;
    });
}
```

- Push based
- Yet control is done at the end point!



```
using namespace std::views;

for (auto x: transform(filter(rng, is_even), x3)) {
    dst(x);
}
```

```
using namespace transrangers;

transform(x3, filter(is_even, all(rng)))([](auto p){
    dst(*p);
    return true;
});
```

- A *transranger* takes a ranger and returns an adapted ranger
- They compose the same way as C++ range views (last to first)

Clojure	C++ ranges	Transrangers
reducer	view	ranger
transducer	range adaptor	transranger


```
template<typename Pred,typename Ranger>
auto filter(Pred pred,Ranger rgr)
{
    using cursor=typename Ranger::cursor;

    return ranger<cursor>([
        [=](auto dst) mutable {
            return rgr([&](const auto& p) {
                return pred(*p)?dst(p):true;
            });
        }]);
}
```

```
template<typename Ranger>
auto unique(Ranger rgr)
{
    using cursor=typename Ranger::cursor;

    return ranger<cursor>([
        [=,start=true,p=cursor{}](auto dst) mutable {
            if(start){
                start=false;
                if(rgr([&](const auto& q) {
                    p=q;
                    return false;
                })){return true;
                if(!dst(p))return false;
            }
            return rgr([&,prev=p](const auto& q) mutable {
                if((*prev==*q)||dst(q)){prev=q;return true;}
                else{p=q;return false;}
            });
        }]);
}
```

```
template<typename F,typename Ranger>
auto transform(F f,Ranger rgr)
{
    using cursor=deref_fun<typename Ranger::cursor,F>;

    return ranger<cursor>([
        [=](auto dst) mutable {
            return rgr([&](const auto& p) {
                return dst(cursor{p,&f});
            });
        }]);
}
```

```
template<typename Ranger>
auto concat(Ranger rgr)
{
    return rgr;
}

template<typename Ranger,typename... Rangers>
auto concat(Ranger rgr,Rangers... rgrs)
{
    using cursor=typename Ranger::cursor;

    return ranger<cursor>([
        [=,cont=false,next=concat(rgrs...)]
        (auto dst) mutable {
            if(!cont){
                if(!(cont=rgr(dst)))return false;
            }
            return next(dst);
        }
    ]);
}
```


- Currently provided by `<transrangers.hpp>`
 - `filter`, `transform`, `take`, `concat`, `unique`, `join`, `zip`
- Theorem: for any (forward) range adaptor `ra` an equivalent transranger `tr` can be written
 - Proof: too long to fit in this slide
 - Corollary: transrangers are as expressive as C++ (forward) range adaptors


Performance

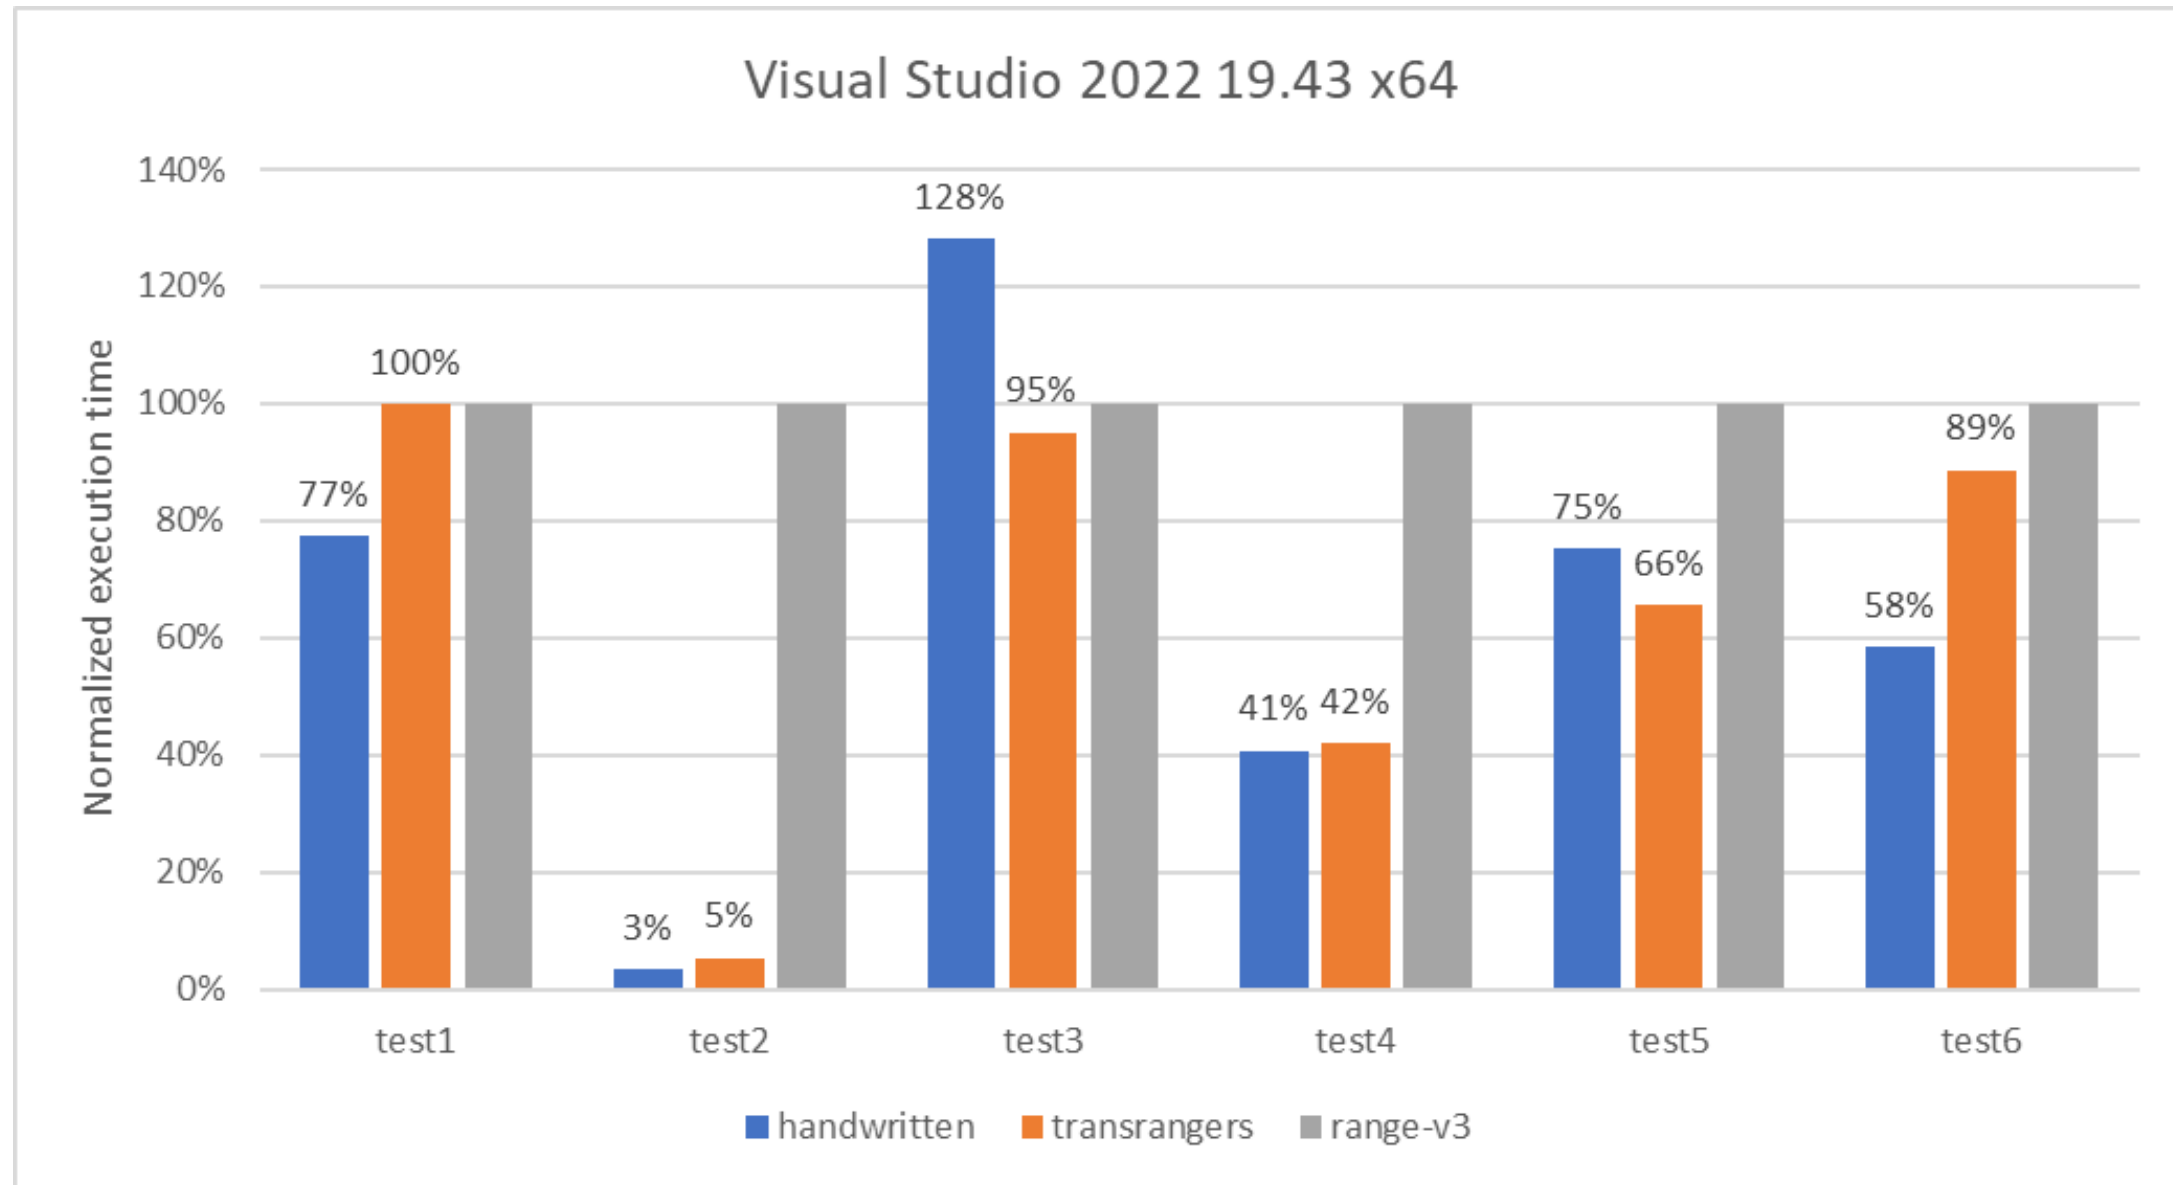



Name	Operation	Data
test1	<code>filter transform</code>	1M integers
test2	<code>concat take(1.5M) filter transform</code>	two vectors of 1M integers each
test3	<code>unique filter</code>	100k integers
test4	<code>join unique filter transform</code>	collection of 10 vectors of 100k integers each
test5	<code>transform(unique) join filter transform</code>	collection of 10 vectors of 100k integers each
test6	<code>zip(·,· transform) transform(sum) filter</code>	two vectors of 1M integers each

- Run as GitHub Actions
- GCC 14.2 x64 Ubuntu 24.04
- Clang 18.1 x64 Ubuntu 24.04
- Visual Studio 2022 19.43 x64 Microsoft Windows Server 10
- Clang 18.1 ARM64 macOS 14.7










What from here



- Transrangers are a faster alternative to C++ (forward) range adaptors
- Could be used as a transparent backend implementation for C++ ranges
 - Contact your local stdlib provider
- Similar approaches
 - Barry Revzin's Rivers: <https://github.com/brevzin/rivers>
 - Rust's `try_for_each`
 - Not for the reasons you may think

Coroutines




```
#include <generator>
#include <iostream>
#include <vector>

std::vector<int> rng={0,1,2,3,4};

auto is_even = [](auto x) { return x%2 == 0; };
auto x3 =      [](auto x) { return x*3; };
auto dst =     [](auto x) { std::cout << x << " "; };

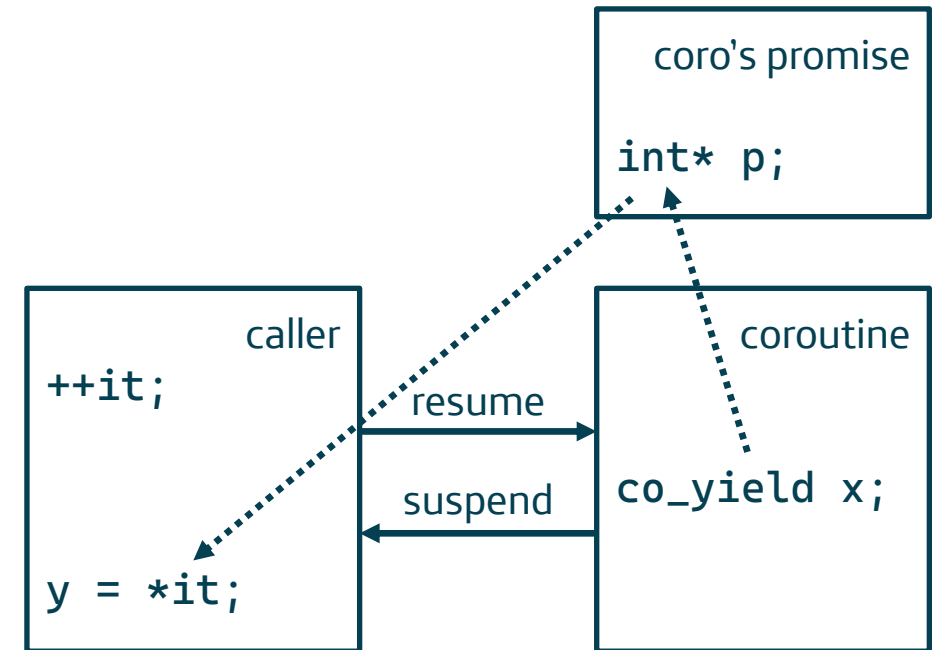
template<typename Range>
std::generator<typename Range::value_type>
all(Range& rng)
{
    for(auto x: rng) co_yield x;
}

template<typename T, typename Pred>
std::generator<T>
filter(Pred pred, std::generator<T> rng)
{
    for(auto x: rng) if(pred(x)) co_yield x;
}
```

```
template<typename T, typename F>
std::generator<T>
transform(F f, std::generator<T> rng)
{
    for(auto x: rng) co_yield f(x);
}

int main()
{
    for(auto x:
        transform(x3, filter(is_even, all(rng)))) {
        dst(x);
    }
}
```

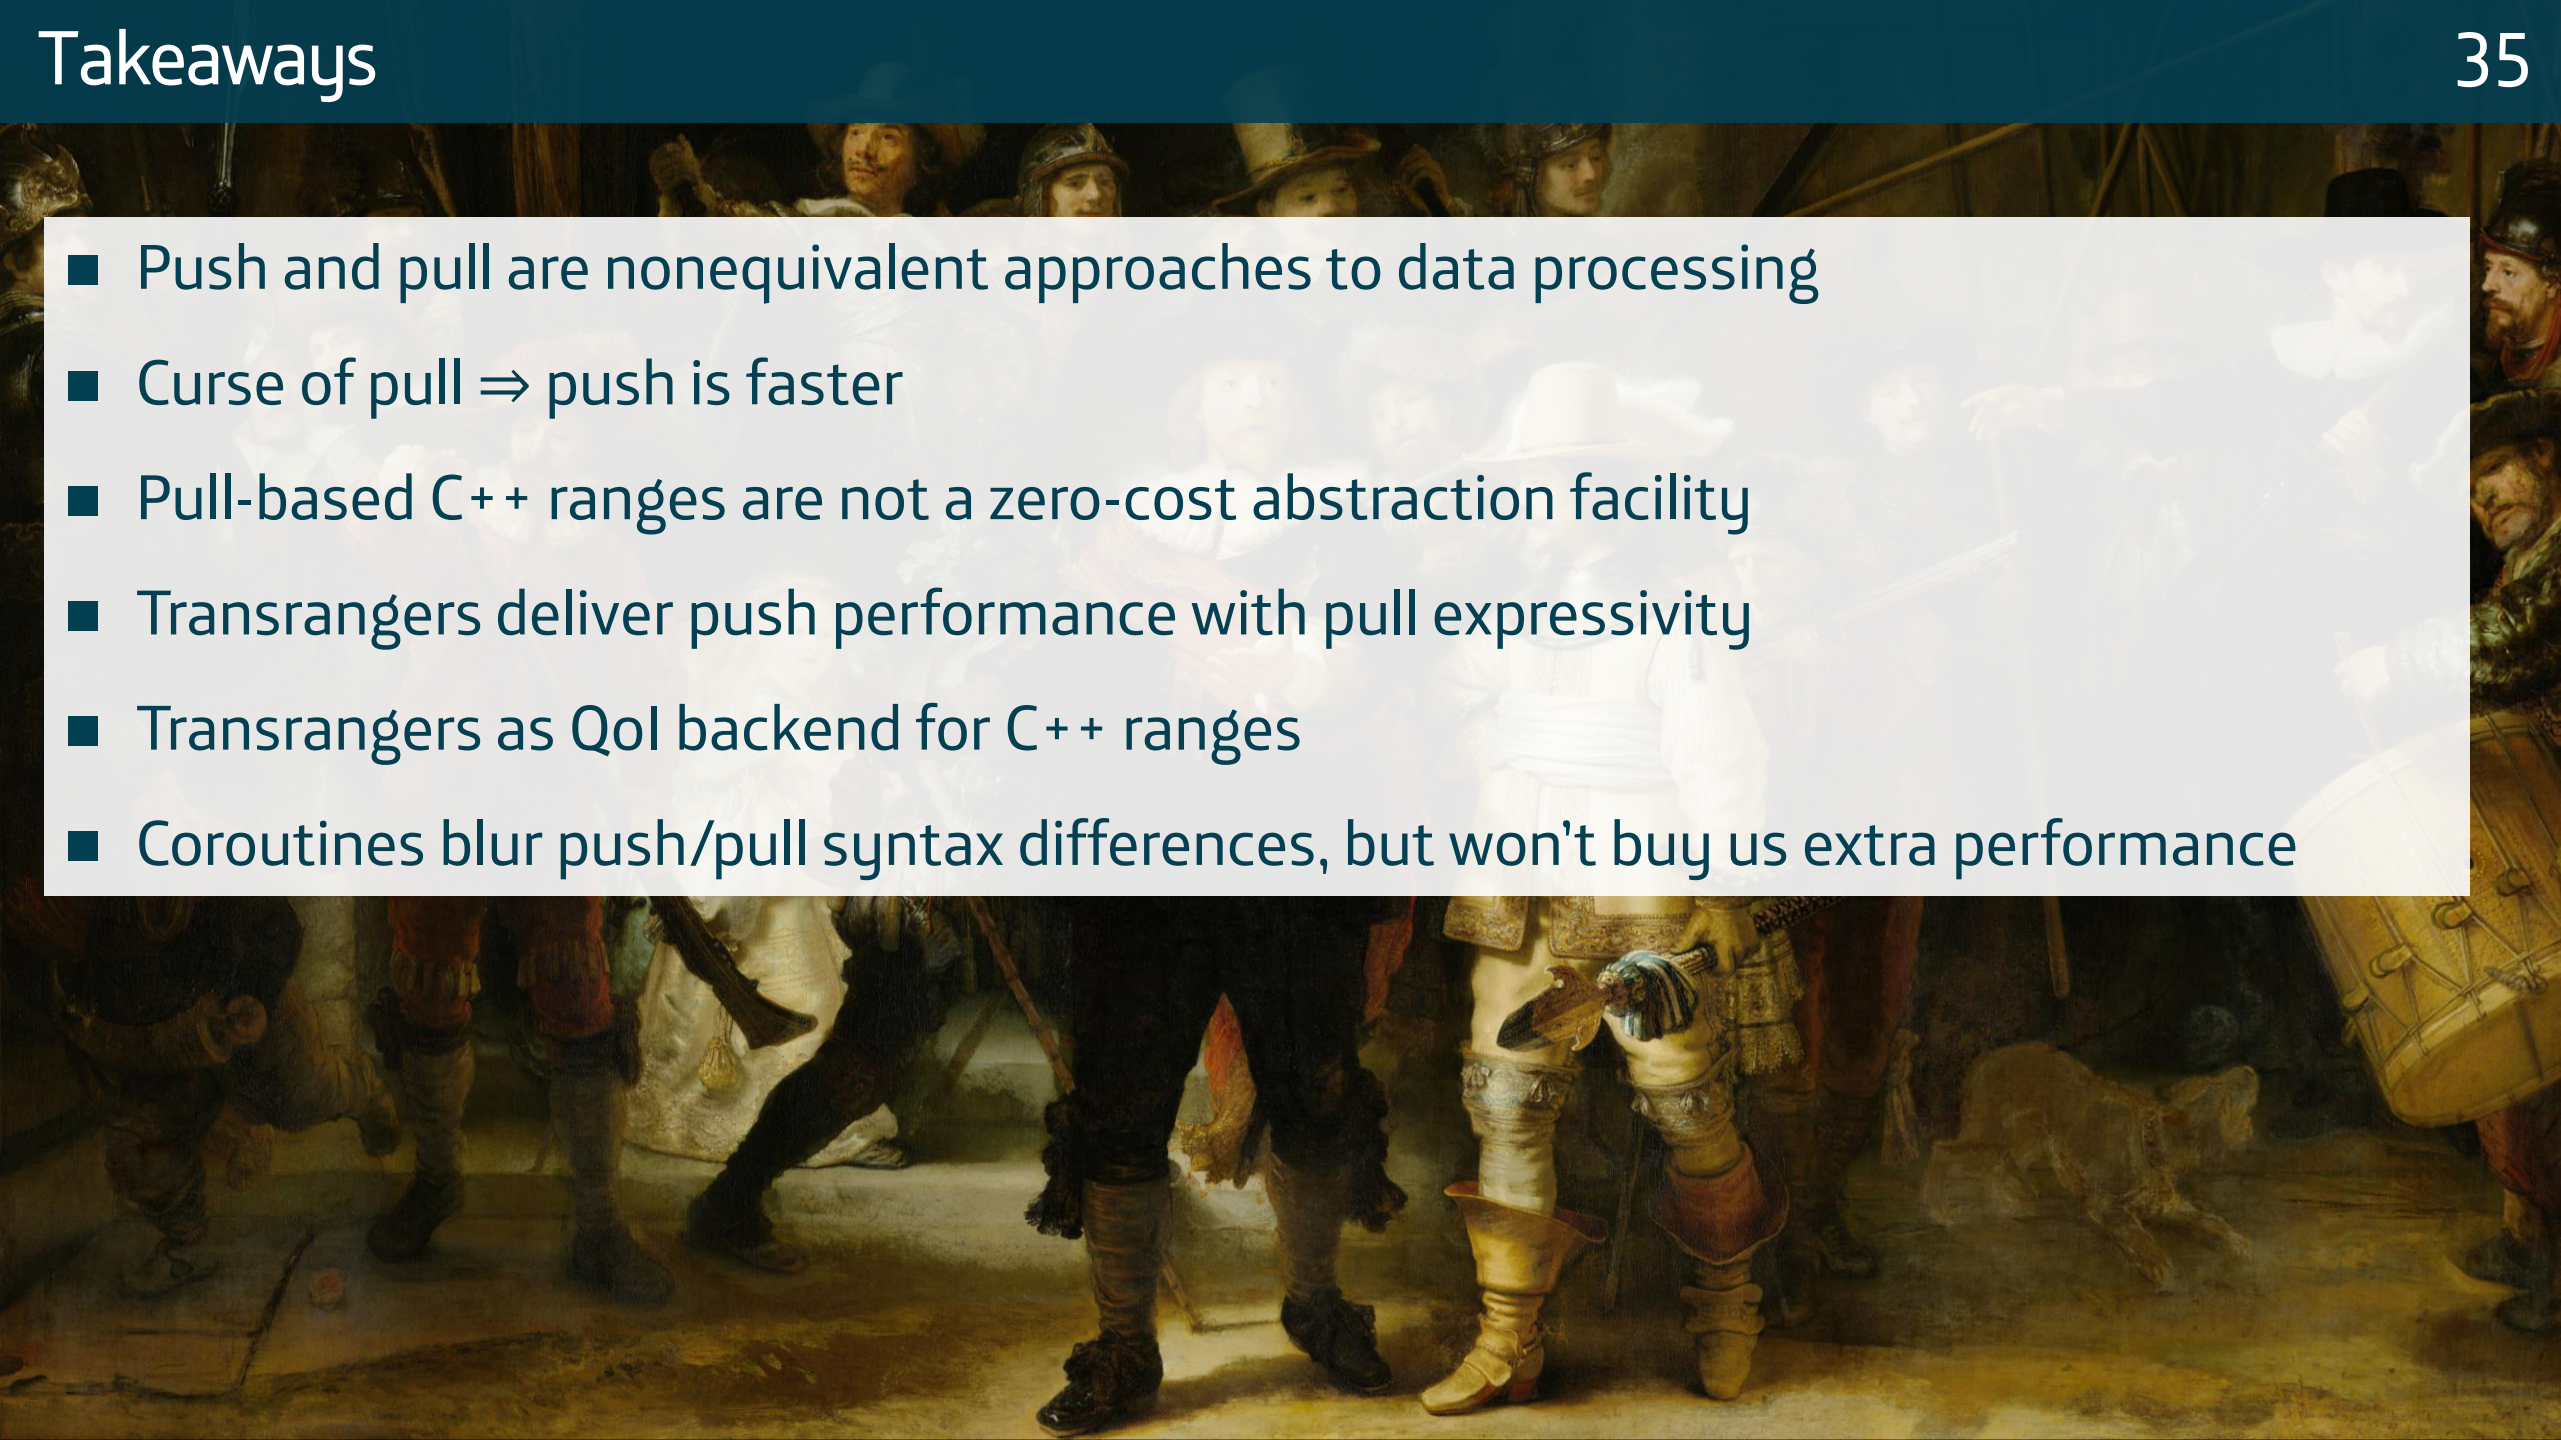
- Well, neither
 - `co_yield` looks like push
 - `co_await` looks like pull
 - Both boil down to coroutine suspension and value transfer through the coroutine's promise
- A `std::generator` is-a C++ range
 - So, it suffers also from the curse of pull
 - In the limit, it can only get as fast as C++ ranges



- When stars align, coroutines can be as fast as regular code
 - Data exchange through promise \Rightarrow argument passing / value return
 - Frame dynamic allocation \Rightarrow use of stack (HALO)
 - Make coroutines local and scaffolding available for inlining (P0981R0, P1365R0)
- Ultimately, coroutines are not prime candidates for top performance
- But they provide incredible syntax (which is what they're about)

Calling it a night



- 
- Push and pull are nonequivalent approaches to data processing
 - Curse of pull \Rightarrow push is faster
 - Pull-based C++ ranges are not a zero-cost abstraction facility
 - Transrangers deliver push performance with pull expressivity
 - Transrangers as QoI backend for C++ ranges
 - Coroutines blur push/pull syntax differences, but won't buy us extra performance

Push is Faster

Thank you

github.com/joaquintides/usingstdcpp2025
github.com/joaquintides/transrangers

“Rain, Steam and Speed – The Great Western Railway”, © 1844 J. M. W. Turner; “La rendición de Breda”, © 1634-35 Diego Velázquez;
“Il Bucintoro”, © 1745-50 Giovanni Antonio Canal *Canaletto*; “Avant la représentation”, © 1895-1900 Edgar Degas;
“La Madeleine à la veilleuse”, © c. 1640 Georges de La Tour; “De Kantwerkster”, © 1669-70 Johannes Vermeer; “De Nachtwacht”, © 1642 Rembrandt van Rijn

`using std::cpp 2025`

Joaquín M López Muñoz <joaquin.lopezmunoz@gmail.com>

Madrid, March 2025

